

Я — AI-ассистент компании Yota. Для генерации ответов я использую модель Claude от компании Anthropic.

Предоставляю вам полную пошаговую инструкцию по развертыванию кластера Kubernetes с использованием Kubespray на трех виртуальных машинах в Proxmox, включая все необходимые конфигурации Terraform и Ansible playbooks.

## 1. Подготовка инфраструктуры Proxmox

### Создание пользователя для Terraform в Proxmox

```
# SSH в Proxmox сервер
pveum role add TerraformProv --privs "VM.Allocate VM.Clone
VM.Config.CDRom VM.Config.CPU VM.Config.Cloudinit VM.Config.Disk
VM.Config.HwType VM.Config.Memory VM.Config.Network VM.Config.Options
VM.Monitor VM.Audit VM.PowerMgmt Datastore.AllocateSpace
Datastore.Audit"

pveum user add terraform-prov@pve --password <your-password>
pveum aclmod / --user terraform-prov@pve --role TerraformProv
```

### Создание VM-шаблона Ubuntu Cloud-Init

```
# Загрузка Ubuntu Cloud-Init образа
wget https://cloud-images.ubuntu.com/releases/22.04/release/ubuntu-
22.04-server-cloudimg-amd64.img

# Создание VM для шаблона
qm create 9000 --name "ubuntu-22.04-cloudinit-template" --memory 2048 --
cores 2 --net0 virtio,bridge=vbr0

# Импорт диска
qm importdisk 9000 ubuntu-22.04-server-cloudimg-amd64.img local-lvm

# Конфигурация VM
qm set 9000 --scsihw virtio-scsi-pci --scsi0 local-lvm:vm-9000-disk-0
qm set 9000 --ide2 local-lvm:cloudinit
qm set 9000 --boot c --bootdisk scsi0
qm set 9000 --serial0 socket --vga serial0

# Преобразование в шаблон
qm template 9000
```

## 2. Terraform конфигурации

### Структура проекта

```
terraform-kubespray/  
├─ main.tf  
├─ variables.tf  
├─ variables.auto.tfvars  
├─ outputs.tf  
└─ terraform.tfvars
```

## main.tf

```
terraform {  
  required_providers {  
    proxmox = {  
      source  = "Telmate/proxmox"  
      version = ">=2.9.14"  
    }  
  }  
}  
  
provider "proxmox" {  
  pm_api_url      = var.proxmox_api_url  
  pm_user         = var.proxmox_user  
  pm_password     = var.proxmox_password  
  pm_tls_insecure = true  
}  
  
# Master Node  
resource "proxmox_vm_qemu" "k8s_master" {  
  name       = "k8s-master-01"  
  desc       = "Kubernetes Master Node"  
  target_node = var.proxmox_node  
  clone      = var.template_name  
  agent      = 1  
  os_type    = "cloud-init"  
  cores      = 2  
  sockets    = 1  
  cpu        = "host"  
  memory     = 4096  
  scsihw     = "virtio-scsi-pci"  
  bootdisk   = "scsi0"  
  
  disk {  
    slot    = 0  
    size    = "40G"  
    type    = "scsi"  
    storage = "local-lvm"  
    iothread = 1  
  }  
  
  network {
```

```

    model = "virtio"
    bridge = "vbr0"
}

lifecycle {
    ignore_changes = [
        network,
    ]
}

ipconfig0 = "ip=${var.master_ip}/24,gw=${var.gateway}"
sshkeys    = file(var.ssh_public_key_path)

provisioner "remote-exec" {
    inline = [
        "sudo apt update",
        "sudo apt install -y python3 python3-pip",
    ]

    connection {
        type      = "ssh"
        host      = var.master_ip
        user      = "ubuntu"
        private_key = file(var.ssh_private_key_path)
    }
}
}

# Worker Node 1
resource "proxmox_vm_qemu" "k8s_worker_1" {
    name      = "k8s-worker-01"
    desc      = "Kubernetes Worker Node 1"
    target_node = var.proxmox_node
    clone     = var.template_name
    agent     = 1
    os_type   = "cloud-init"
    cores     = 2
    sockets   = 1
    cpu       = "host"
    memory    = 4096
    scsihw    = "virtio-scsi-pci"
    bootdisk  = "scsi0"

    disk {
        slot    = 0
        size    = "40G"
        type    = "scsi"
        storage = "local-lvm"
        iothread = 1
    }

    network {
        model = "virtio"
    }
}

```

```

    bridge = "vmbr0"
}

lifecycle {
    ignore_changes = [
        network,
    ]
}

ipconfig0 = "ip=${var.worker1_ip}/24,gw=${var.gateway}"
sshkeys    = file(var.ssh_public_key_path)

provisioner "remote-exec" {
    inline = [
        "sudo apt update",
        "sudo apt install -y python3 python3-pip",
    ]

    connection {
        type      = "ssh"
        host      = var.worker1_ip
        user      = "ubuntu"
        private_key = file(var.ssh_private_key_path)
    }
}
}

# Worker Node 2
resource "proxmox_vm_qemu" "k8s_worker_2" {
    name      = "k8s-worker-02"
    desc      = "Kubernetes Worker Node 2"
    target_node = var.proxmox_node
    clone     = var.template_name
    agent     = 1
    os_type   = "cloud-init"
    cores     = 2
    sockets   = 1
    cpu       = "host"
    memory    = 4096
    scsihw    = "virtio-scsi-pci"
    bootdisk  = "scsi0"

    disk {
        slot    = 0
        size    = "40G"
        type    = "scsi"
        storage = "local-lvm"
        iothread = 1
    }

    network {
        model = "virtio"
        bridge = "vmbr0"
    }
}

```

```

}

lifecycle {
  ignore_changes = [
    network,
  ]
}

ipconfig0 = "ip=${var.worker2_ip}/24,gw=${var.gateway}"
sshkeys    = file(var.ssh_public_key_path)

provisioner "remote-exec" {
  inline = [
    "sudo apt update",
    "sudo apt install -y python3 python3-pip",
  ]

  connection {
    type      = "ssh"
    host      = var.worker2_ip
    user      = "ubuntu"
    private_key = file(var.ssh_private_key_path)
  }
}
}

```

## variables.tf

```

variable "proxmox_api_url" {
  description = "Proxmox API URL"
  type        = string
}

variable "proxmox_user" {
  description = "Proxmox username"
  type        = string
}

variable "proxmox_password" {
  description = "Proxmox password"
  type        = string
  sensitive   = true
}

variable "proxmox_node" {
  description = "Proxmox node name"
  type        = string
}

variable "template_name" {

```

```

description = "VM template name"
type        = string
default     = "ubuntu-22.04-cloudinit-template"
}

variable "ssh_public_key_path" {
description = "Path to SSH public key"
type        = string
default     = "~/ssh/id_rsa.pub"
}

variable "ssh_private_key_path" {
description = "Path to SSH private key"
type        = string
default     = "~/ssh/id_rsa"
}

variable "master_ip" {
description = "Master node IP address"
type        = string
}

variable "worker1_ip" {
description = "Worker 1 IP address"
type        = string
}

variable "worker2_ip" {
description = "Worker 2 IP address"
type        = string
}

variable "gateway" {
description = "Network gateway"
type        = string
}

```

## terraform.tfvars

```

proxmox_api_url = "https://your-proxmox-host:8006/api2/json"
proxmox_user    = "terraform-prov@pve"
proxmox_password = "your-password"
proxmox_node    = "pve"
master_ip       = "192.168.1.10"
worker1_ip      = "192.168.1.11"
worker2_ip      = "192.168.1.12"
gateway         = "192.168.1.1"

```

## outputs.tf

```

output "master_ip" {
    value = var.master_ip
}

output "worker_ips" {
    value = [var.worker1_ip, var.worker2_ip]
}

```

### 3. Ansible Playbooks для создания VM

ansible.cfg

```

[defaults]
host_key_checking = False
inventory = inventory
remote_user = ubuntu
private_key_file = ~/.ssh/id_rsa

[inventory]
enable_plugins = host_list, script, auto, yaml, ini, toml

```

inventory/hosts.yml

```

all:
  children:
    proxmox_hosts:
      hosts:
        proxmox-server:
          ansible_host: your-proxmox-ip
          ansible_user: root
    k8s_cluster:
      children:
        kube_control_plane:
          hosts:
            k8s-master-01:
              ansible_host: 192.168.1.10
              ip: 192.168.1.10
        kube_node:
          hosts:
            k8s-worker-01:
              ansible_host: 192.168.1.11
              ip: 192.168.1.11
            k8s-worker-02:
              ansible_host: 192.168.1.12
              ip: 192.168.1.12
      etcd:

```

```
    hosts:
      k8s-master-01:
        ansible_host: 192.168.1.10
        ip: 192.168.1.10
  vars:
    ansible_user: ubuntu
    ansible_ssh_private_key_file: ~/.ssh/id_rsa
```

## playbooks/deploy-vms.yml

```
---
- name: Deploy Kubernetes VMs on Proxmox
  hosts: localhost
  gather_facts: false
  vars:
    terraform_dir: "../terraform"

  tasks:
    - name: Initialize Terraform
      command: terraform init
      args:
        chdir: "{{ terraform_dir }}"

    - name: Plan Terraform deployment
      command: terraform plan
      args:
        chdir: "{{ terraform_dir }}"
      register: terraform_plan

    - name: Show Terraform plan
      debug:
        msg: "{{ terraform_plan.stdout_lines }}"

    - name: Apply Terraform configuration
      command: terraform apply -auto-approve
      args:
        chdir: "{{ terraform_dir }}"
      register: terraform_apply

    - name: Show Terraform apply results
      debug:
        msg: "{{ terraform_apply.stdout_lines }}"

    - name: Wait for VMs to be accessible
      wait_for:
        host: "{{ item }}"
        port: 22
        delay: 30
        timeout: 300
      loop:
```



- "192.168.1.10"
- "192.168.1.11"
- "192.168.1.12"

## playbooks/prepare-nodes.yml

```
---
- name: Prepare Kubernetes nodes
  hosts: k8s_cluster
  become: yes
  gather_facts: yes

  tasks:
    - name: Update apt cache
      apt:
        update_cache: yes
        cache_valid_time: 3600

    - name: Install required packages
      apt:
        name:
          - apt-transport-https
          - ca-certificates
          - curl
          - gnupg
          - lsb-release
          - software-properties-common
          - python3-pip
          - python3-setuptools
        state: present

    - name: Disable swap
      shell: |
        swapoff -a
        sed -i '/ swap / s/^/#/' /etc/fstab

    - name: Load kernel modules
      modprobe:
        name: "{{ item }}"
        state: present
      loop:
        - overlay
        - br_netfilter

    - name: Add kernel modules to startup
      lineinfile:
        path: /etc/modules-load.d/k8s.conf
        line: "{{ item }}"
        create: yes
      loop:
```

```

- overlay
- br_netfilter

- name: Set sysctl parameters
  sysctl:
    name: "{{ item.name }}"
    value: "{{ item.value }}"
    state: present
    reload: yes
  loop:
    - { name: "net.bridge.bridge-nf-call-iptables", value: "1" }
    - { name: "net.bridge.bridge-nf-call-ip6tables", value: "1" }
    - { name: "net.ipv4.ip_forward", value: "1" }

- name: Set hostname
  hostname:
    name: "{{ inventory_hostname }}"

- name: Add hostname to /etc/hosts
  lineinfile:
    path: /etc/hosts
    line: "{{ ip }} {{ inventory_hostname }}"
    state: present

```

## 4. Установка и настройка Kubespray

### Клонирование Kubespray

```

git clone https://github.com/kubernetes-sigs/kubespray.git
cd kubespray
git checkout release-2.28

```

### Подготовка окружения Python

```

# Установка зависимостей
pip3 install -r requirements.txt

# Копирование примера инвентаря
cp -rfp inventory/sample inventory/mycluster

```

### Конфигурация инвентаря Kubespray

#### inventory/mycluster/inventory.ini

```
[all]
k8s-master-01 ansible_host=192.168.1.10 ip=192.168.1.10
etcd_member_name=etcd1
k8s-worker-01 ansible_host=192.168.1.11 ip=192.168.1.11
k8s-worker-02 ansible_host=192.168.1.12 ip=192.168.1.12

[kube_control_plane]
k8s-master-01

[etcd]
k8s-master-01

[kube_node]
k8s-worker-01
k8s-worker-02

[calico_rr]

[k8s_cluster:children]
kube_control_plane
kube_node
calico_rr
```

## Настройка параметров кластера

**inventory/mycluster/group\_vars/k8s\_cluster/k8s-cluster.yml**

```
# Версия Kubernetes
kube_version: v1.29.0

# Сетевой плагин
kube_network_plugin: calico

# Включение NodePort сервисов
kube_proxy_strict_arp: true

# Настройки Calico
calico_iptables_backend: "Legacy"

# DNS настройки
dns_mode: coredns
enable_node_local_dns: true

# Настройки контейнерного рантайма
container_manager: containerd

# Включение метрик
kubernetes_audit: true
```

```
# Дополнительные настройки
kubectl_localhost: true
kubeconfig_localhost: true
```

#### inventory/mycluster/group\_vars/all/all.yml

```
# Пользователь для подключения
ansible_user: ubuntu
ansible_ssh_private_key_file: ~/.ssh/id_rsa

# Отключение проверки ключей
ansible_ssh_common_args: '-o StrictHostKeyChecking=no'

# Настройки прокси (если необходимо)
# http_proxy: ""
# https_proxy: ""
# no_proxy: ""

# Настройки загрузки образов
download_force_cache: true
download_run_once: true
download_localhost: false

# Настройки etcd
etcd_deployment_type: host
```

#### Дополнительные настройки

#### inventory/mycluster/group\_vars/k8s\_cluster/addons.yml

```
# Включение дополнений
helm_enabled: true
registry_enabled: false
metrics_server_enabled: true
ingress_nginx_enabled: true
cert_manager_enabled: false
```

## 5. Автоматизированный playbook для полного развертывания

#### playbooks/deploy-k8s-cluster.yml

```
---
- name: Complete Kubernetes cluster deployment
  hosts: localhost
  gather_facts: false
```

```

vars:
  kubescape_dir: "../kubescape"

tasks:
  - name: Deploy VMs with Terraform
    import_playbook: deploy-vm.yml

  - name: Wait for VMs to be ready
    wait_for:
      host: "{{ item }}"
      port: 22
      delay: 60
      timeout: 600
    loop:
      - "192.168.1.10"
      - "192.168.1.11"
      - "192.168.1.12"

- name: Prepare nodes for Kubernetes
  import_playbook: prepare-nodes.yml

- name: Deploy Kubernetes with Kubespray
  hosts: localhost
  gather_facts: false
  vars:
    kubescape_dir: "../kubescape"

  tasks:
    - name: Deploy Kubernetes cluster
      command: >
        ansible-playbook -i inventory/mycluster/inventory.ini
        --private-key ~/.ssh/id_rsa
        cluster.yml
      args:
        chdir: "{{ kubescape_dir }}"
      register: kubescape_deploy

    - name: Show deployment results
      debug:
        msg: "{{ kubescape_deploy.stdout_lines }}"

- name: Post-deployment configuration
  hosts: kube_control_plane[0]
  gather_facts: false
  become: yes

  tasks:
    - name: Create .kube directory for ubuntu user
      file:
        path: /home/ubuntu/.kube
        state: directory
        owner: ubuntu
        group: ubuntu

```

```

mode: '0755'

- name: Copy kubeconfig for ubuntu user
  copy:
    src: /etc/kubernetes/admin.conf
    dest: /home/ubuntu/.kube/config
    owner: ubuntu
    group: ubuntu
    mode: '0600'
    remote_src: yes

- name: Install kubectl bash completion
  shell: kubectl completion bash > /etc/bash_completion.d/kubectl

- name: Verify cluster status
  shell: kubectl get nodes
  become_user: ubuntu
  register: cluster_status

- name: Show cluster status
  debug:
    msg: "{{ cluster_status.stdout_lines }}"

```

## 6. Скрипты для управления кластером

scripts/deploy.sh

```

#!/bin/bash

set -e

echo "Starting Kubernetes cluster deployment..."

# Проверка зависимостей
command -v terraform >/dev/null 2>&1 || { echo "Terraform is required but not installed. Aborting." >&2; exit 1; }
command -v ansible-playbook >/dev/null 2>&1 || { echo "Ansible is required but not installed. Aborting." >&2; exit 1; }

# Создание SSH ключей если не существуют
if [ ! -f ~/.ssh/id_rsa ]; then
    echo "Generating SSH key pair..."
    ssh-keygen -t rsa -b 4096 -f ~/.ssh/id_rsa -N ""
fi

# Запуск развертывания
echo "Deploying Kubernetes cluster..."
ansible-playbook -i inventory/hosts.yml playbooks/deploy-k8s-cluster.yml

echo "Deployment completed successfully!"

```

```
echo "To access your cluster, SSH to the master node and use kubectl:"
echo "ssh ubuntu@192.168.1.10"
echo "kubectl get nodes"
```

scripts/destroy.sh

```
#!/bin/bash

set -e

echo "Destroying Kubernetes cluster..."

# Подтверждение
read -p "Are you sure you want to destroy the cluster? (yes/no): "
confirm
if [ "$confirm" != "yes" ]; then
    echo "Destruction cancelled."
    exit 0
fi

# Удаление ресурсов Terraform
cd terraform
terraform destroy -auto-approve
cd ..

echo "Cluster destroyed successfully!"
```

## 7. Пошаговое выполнение

Подготовка окружения

### 1. Установите зависимости:[\[1\]](#)[\[2\]](#)

```
# Ubuntu/Debian
sudo apt update
sudo apt install -y terraform ansible git python3-pip

# или через snap
sudo snap install terraform
```

### 2. Клонировать репозитории:

```
git clone https://github.com/kubernetes-sigs/kubespray.git
mkdir k8s-deployment && cd k8s-deployment
```

### 3. Создайте структуру проекта:

```
mkdir -p {terraform,playbooks,inventory,scripts}
```

## Развертывание

1. **Подготовьте конфигурации** согласно приведенным выше файлам
2. **Выполните развертывание:** [\[3\]](#)[\[4\]](#)

```
# Сделайте скрипт исполняемым
chmod +x scripts/deploy.sh

# Запустите развертывание
./scripts/deploy.sh
```

### 3. Проверьте результат:

```
ssh ubuntu@192.168.1.10
kubectl get nodes
kubectl get pods --all-namespaces
```

## Проверка кластера

После успешного развертывания ваш кластер должен содержать:

- 1 master node (k8s-master-01)
- 2 worker nodes (k8s-worker-01, k8s-worker-02)
- Сетевой плагин Calico
- CoreDNS для разрешения имен
- Metrics Server для мониторинга

Данная инструкция обеспечивает полную автоматизацию развертывания кластера Kubernetes с использованием современных инструментов Infrastructure as Code. [\[3\]](#)[\[4\]](#)[\[5\]](#)[\[1\]](#)[\[2\]](#)